

## Lesson 22: Regularization

Data Science with Python – BSc Course

45 Minutes

**The Problem:** With many predictors, our model can memorize noise instead of learning patterns. How do we build models that generalize to unseen data?

**After this lesson, you will be able to:**

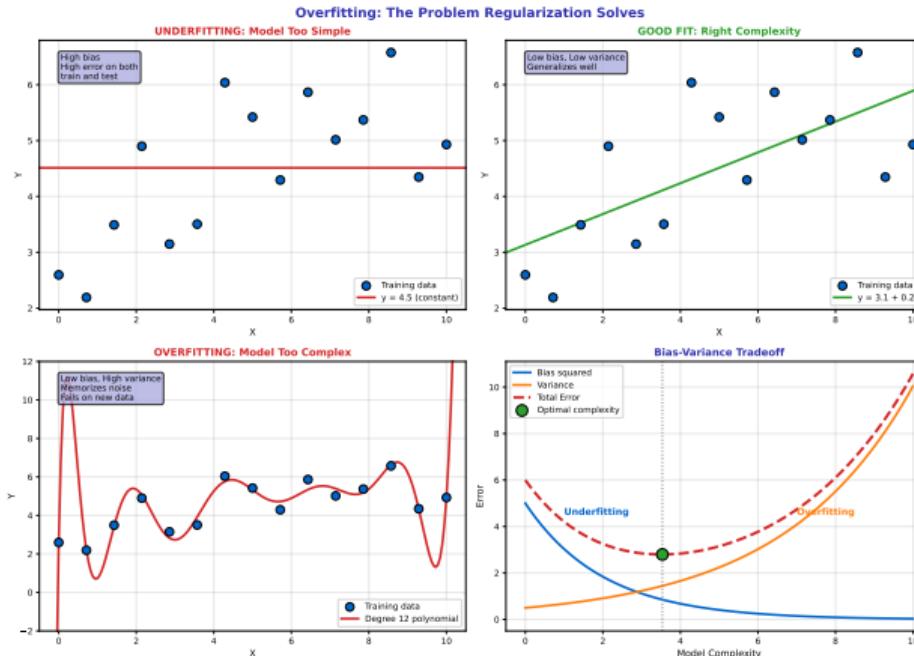
- Recognize overfitting and its causes
- Apply Ridge (L2) regularization to shrink coefficients
- Apply Lasso (L1) for automatic feature selection
- Tune the regularization strength with cross-validation

**Finance Application: Building robust factor models with many correlated predictors**

# The Overfitting Problem

## When Models Memorize Instead of Learn

- High training accuracy but poor test performance
- Complex models fit noise in the training data

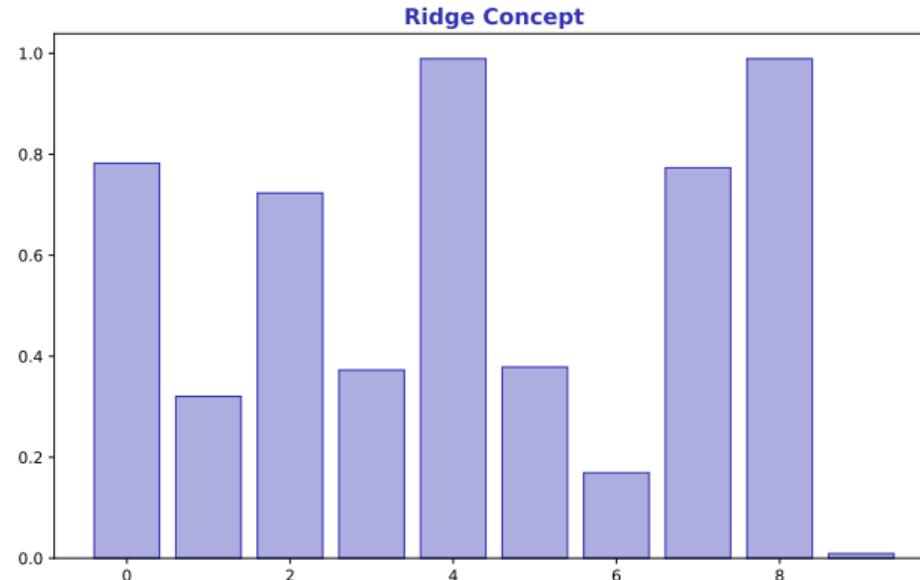


Red flag: If train error keeps dropping but test error rises, you're overfitting

# Ridge Regression (L2)

## Shrink All Coefficients Toward Zero

- Add penalty:  $\text{Loss} = \sum(y - \hat{y})^2 + \lambda \sum \beta_j^2$
- Large  $\lambda$  = stronger shrinkage, simpler model

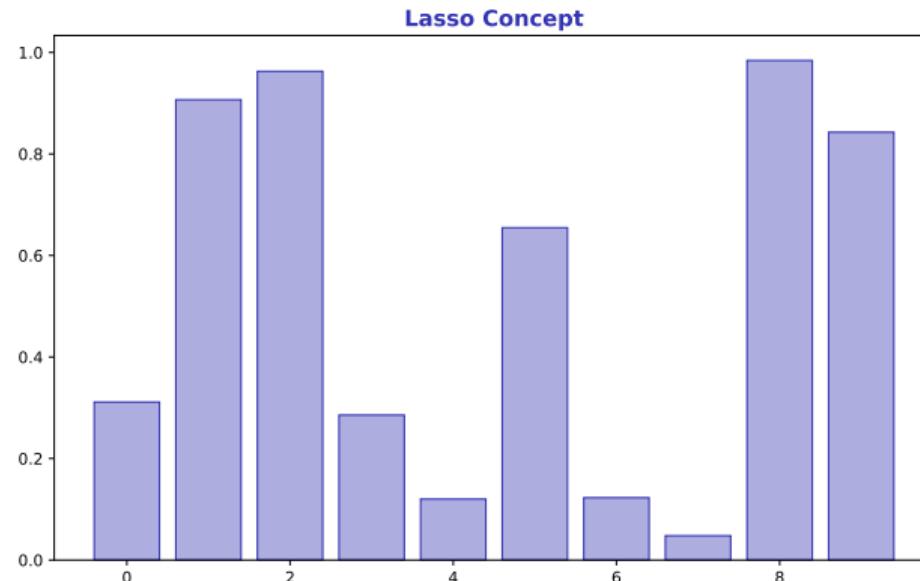


Ridge keeps all features but reduces their influence – good for multicollinearity

# Lasso Regression (L1)

## Some Coefficients Go to Exactly Zero

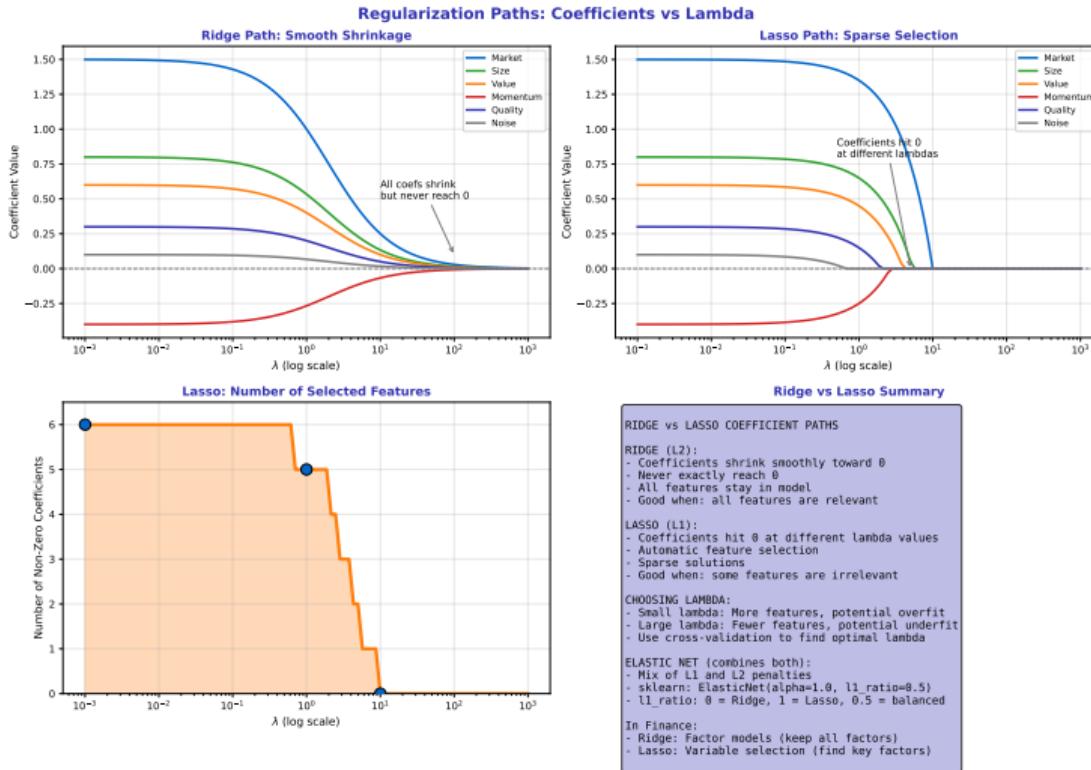
- Add penalty:  $\text{Loss} = \sum(y - \hat{y})^2 + \lambda \sum |\beta_j|$
- L1 penalty creates sparse solutions (automatic feature selection)



Lasso eliminates irrelevant features – use when you suspect many predictors are useless

# Coefficient Paths

## How Coefficients Change with Lambda

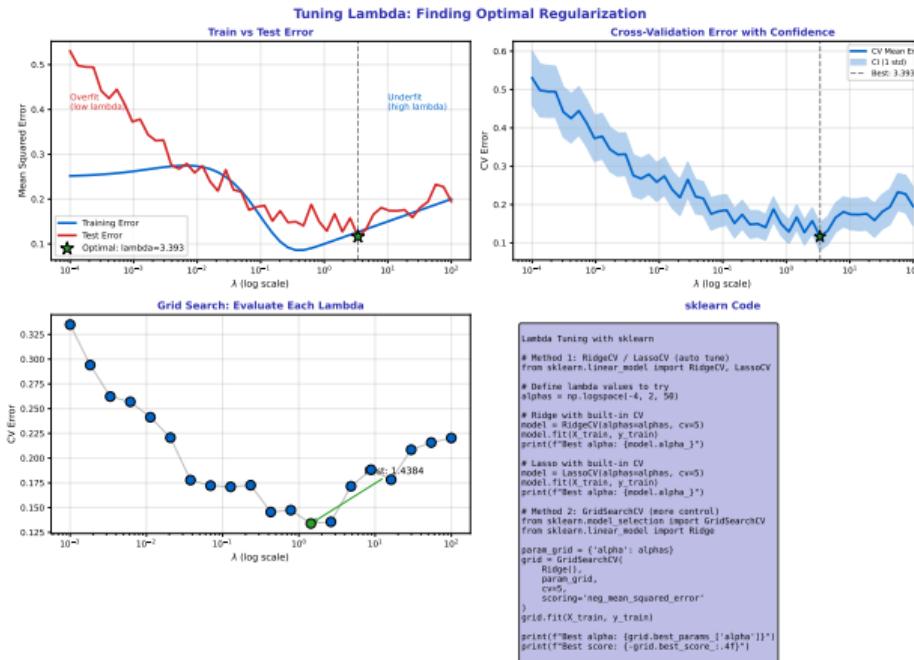


As  $\lambda$  increases, coefficients shrink. Lasso drives some to zero. Ridge does not.

# Tuning Lambda

## Finding the Right Penalty Strength

- Too small: overfitting (model too complex)
- Too large: underfitting (model too simple)

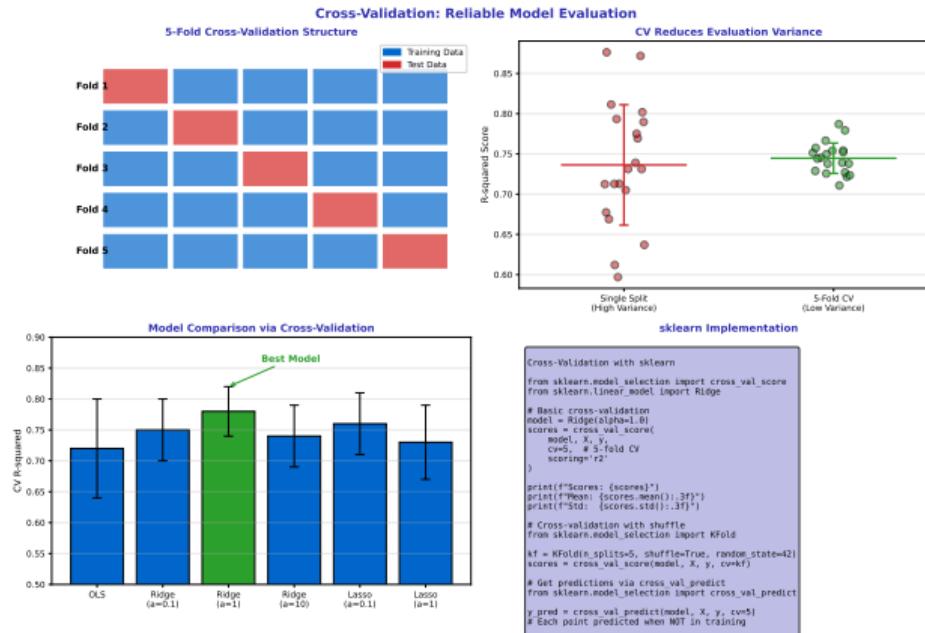


Use cross-validation to find the lambda that minimizes test error

# Cross-Validation for Lambda

## sklearn Makes It Easy

- RidgeCV and LassoCV automatically search lambda values
- K-fold CV: split data K ways, train on K-1, test on 1, average

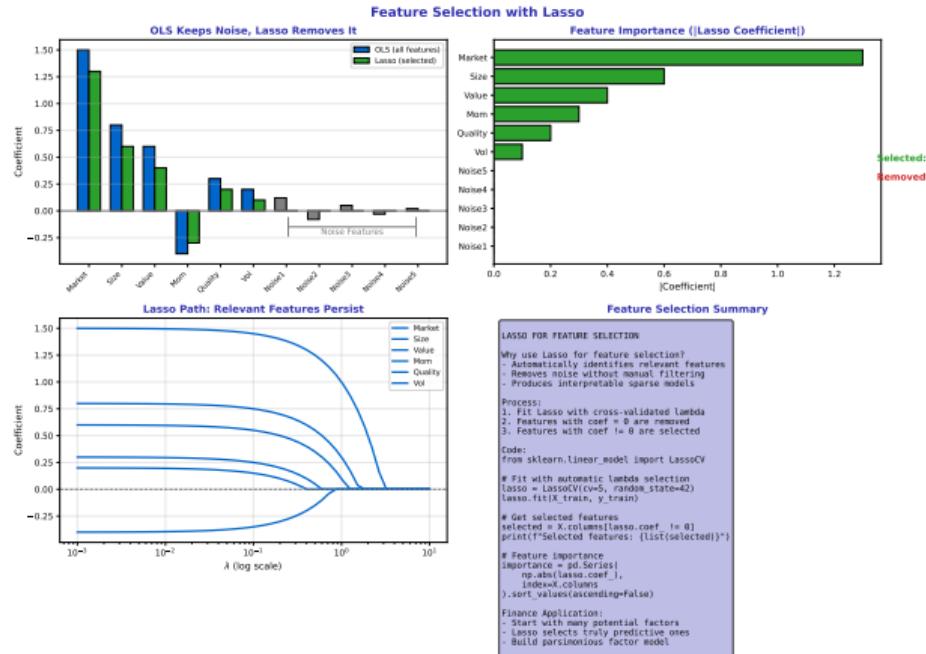


Rule: Use `RidgeCV(alphas=[0.1, 1, 10, 100])` to search logarithmically

# Feature Selection with Lasso

## Which Predictors Actually Matter?

- Non-zero Lasso coefficients = selected features
- Zero coefficients = features eliminated by the model

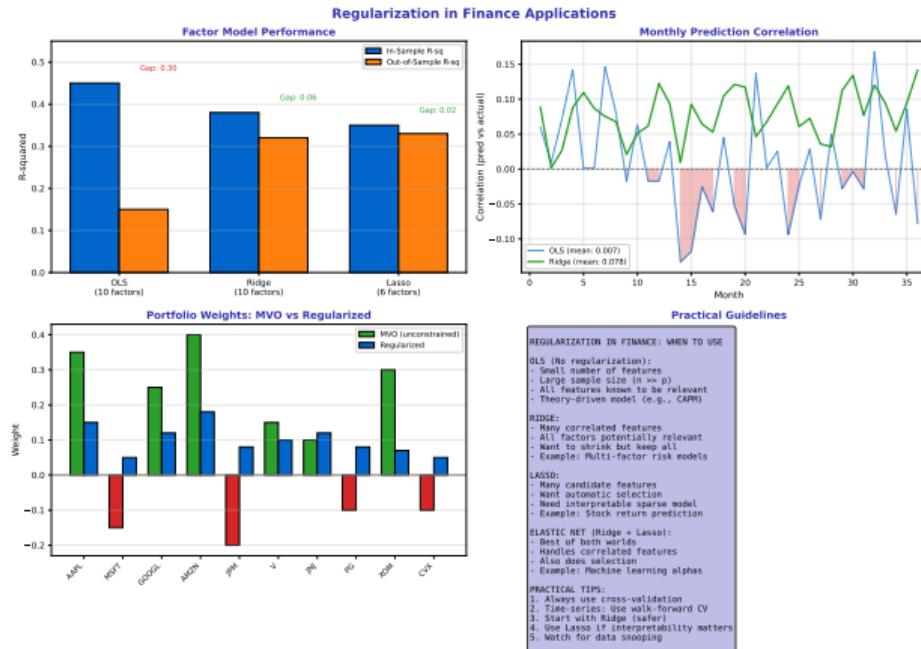


Finance insight: Lasso often keeps 3-5 factors from a candidate set of 20+

# Finance Application: Factor Models

## Building Robust Return Predictions

- Many candidate factors are correlated (value, quality, momentum...)
- Regularization prevents unstable, extreme factor weights



Industry practice: Regularized regression for combining alpha signals

## Hands-On Exercise (25 min)

### Task: Compare Ridge vs Lasso on Multi-Factor Data

- ① Create synthetic data with 20 features (only 5 are truly predictive)
- ② Fit OLS, Ridge, and Lasso models
- ③ Compare test set  $R^2$  for each model
- ④ Plot Lasso coefficients – which features were selected?

**Deliverable:** Bar chart of coefficients comparing OLS vs Ridge vs Lasso.

**Extension:** Use LassoCV to find optimal lambda and report selected features

## Lesson Summary

**Problem Solved:** Regularization prevents overfitting when we have many predictors or limited data.

**Key Takeaways:**

- Ridge (L2) shrinks all coefficients – handles multicollinearity
- Lasso (L1) sets some coefficients to zero – automatic feature selection
- Use cross-validation (RidgeCV, LassoCV) to tune  $\lambda$

**Next Lesson:** Regression Metrics (L23) – how do we measure model quality?

**Memory:** Ridge = Ridge keeps all features. Lasso = Lasso Loses features (L for Lose).