

Lesson 33: Perceptron

Data Science with Python – BSc Course

45 Minutes

The Problem: Traditional ML models have fixed architectures. How do neural networks learn? Understanding the perceptron is the first step toward deep learning.

After this lesson, you will be able to:

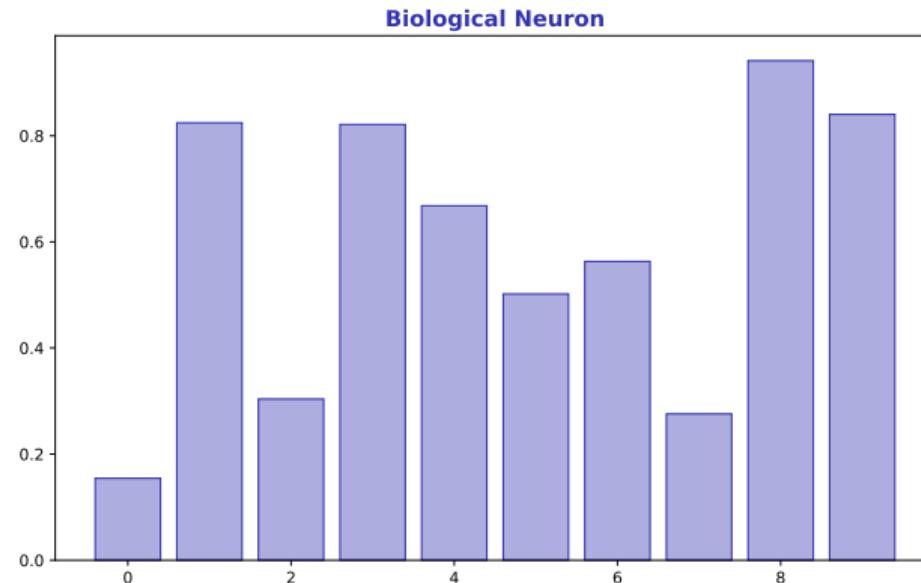
- Understand the biological inspiration for neural networks
- Build and train a single perceptron
- Recognize the limitation: only linearly separable problems
- Prepare for multi-layer networks that overcome this

Finance Application: Foundation for deep learning models in quantitative finance

Biological Neuron

Inspiration from the Brain

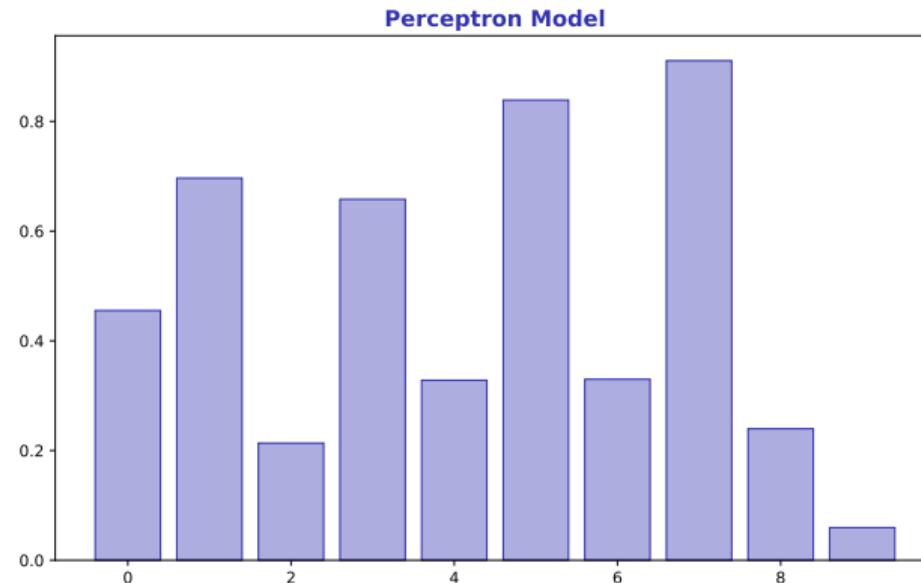
- Dendrites receive signals, soma processes, axon outputs
- Neuron “fires” when input exceeds threshold



Perceptron: simplified mathematical model of biological neuron

The Simplest Neural Network

- Output: $y = \sigma(\sum w_i x_i + b)$ where σ is activation
- Weights w_i control importance of each input

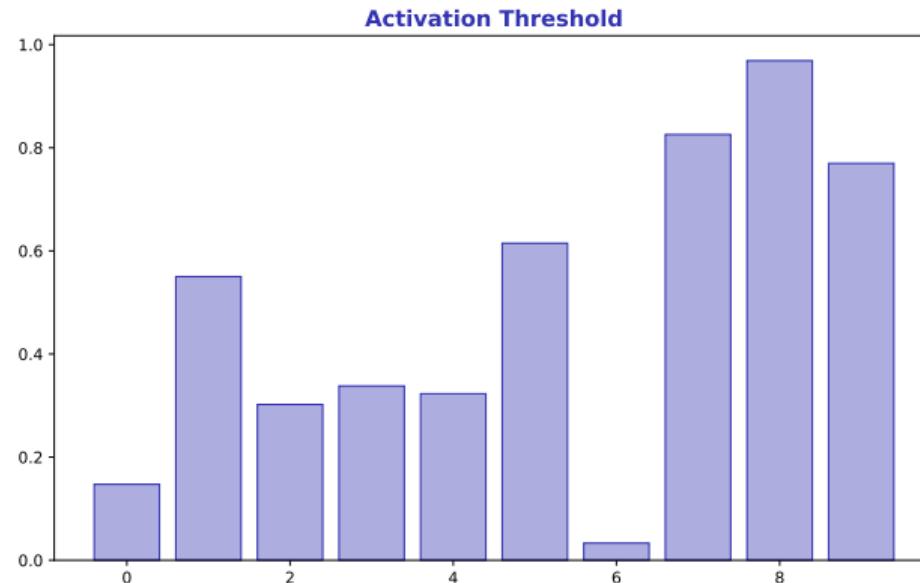


Single perceptron = logistic regression (same math, different framing)

Activation Threshold

When Does the Neuron Fire?

- Step function: output 1 if weighted sum > 0 , else 0
- Modern: sigmoid, ReLU for smooth gradients

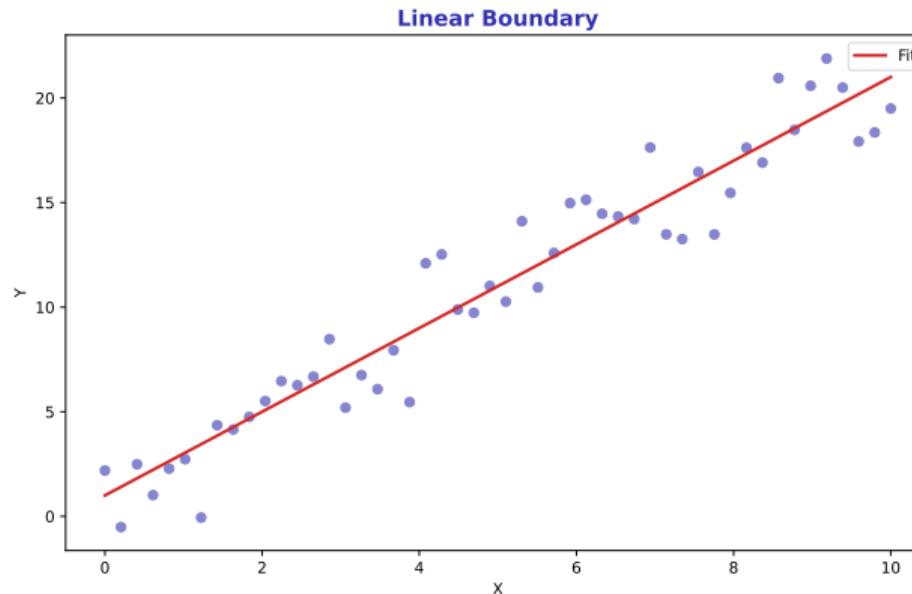


Step function: original perceptron. Sigmoid: smooth for gradient descent.

Linear Decision Boundary

What Can a Perceptron Learn?

- Single perceptron creates a linear boundary (hyperplane)
- Can separate AND, OR but not more complex patterns

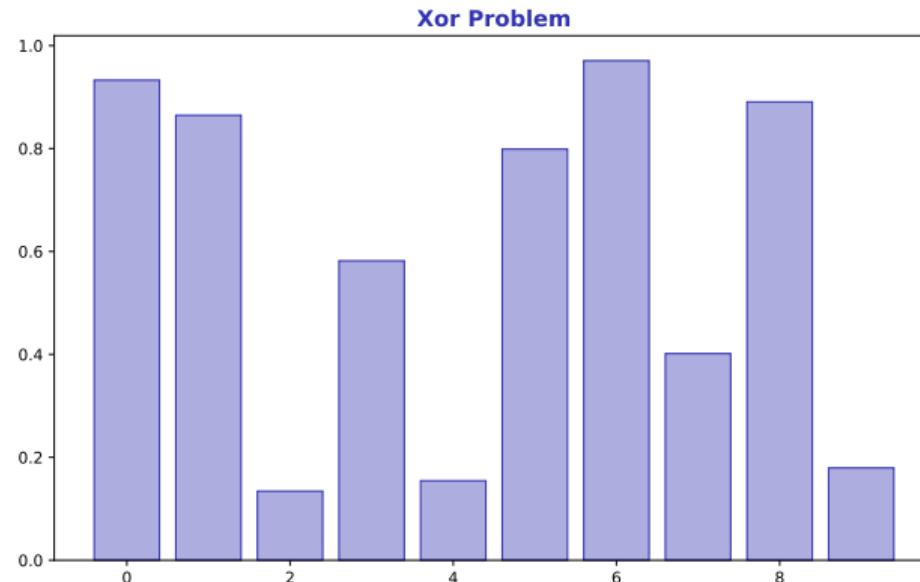


Perceptron = linear classifier. Same limitation as logistic regression.

The XOR Problem

A Famous Limitation

- XOR is not linearly separable – no single line can divide it
- This limitation stalled neural network research for years

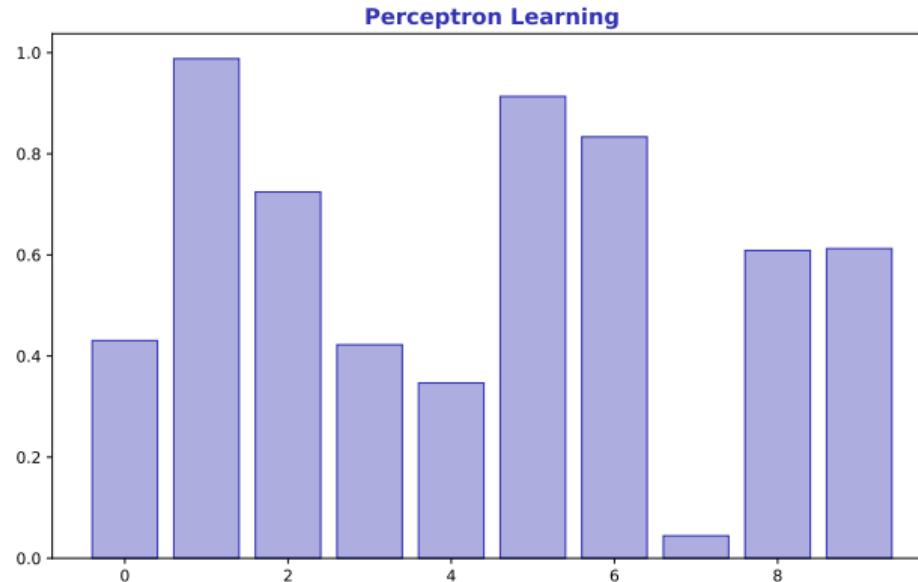


Solution: add hidden layers (multi-layer perceptron) to learn non-linear boundaries

Perceptron Learning Rule

How Weights Are Updated

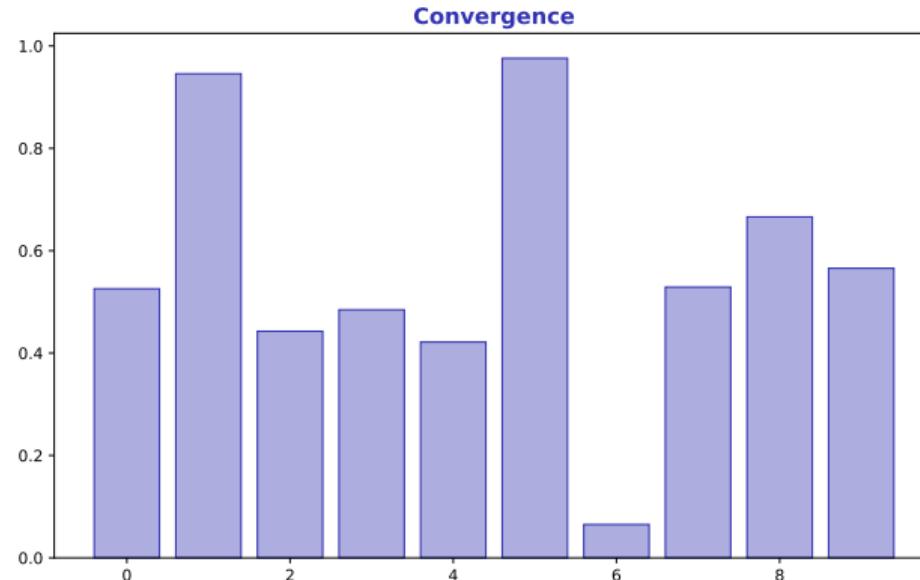
- If wrong: $w_{new} = w_{old} + \eta \cdot (y_{true} - y_{pred}) \cdot x$
- Learning rate η controls step size



Perceptron learning rule: simple but powerful for linearly separable data

When Will Training Stop?

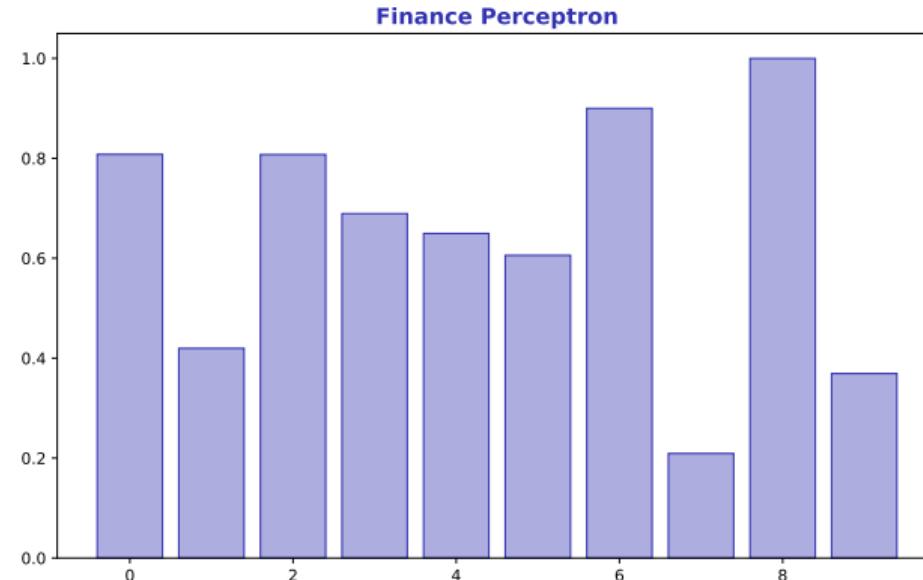
- If data is linearly separable, perceptron converges in finite steps
- If not separable, it oscillates forever (need multi-layer)



Perceptron Convergence Theorem (1962): guaranteed for separable problems

Simple Trading Signals

- Inputs: momentum, volatility, volume signals
- Output: buy (1) or sell (0) decision



Reality: Financial patterns are rarely linearly separable – need deeper networks

Hands-On Exercise (25 min)

Task: Implement Perceptron from Scratch

- ① Create synthetic linearly separable 2D data
- ② Implement perceptron learning rule in Python (no sklearn)
- ③ Train and plot decision boundary at each epoch
- ④ Try on XOR data – observe failure to converge
- ⑤ Compare to sklearn's Perceptron class

Deliverable: Animation of boundary evolution + XOR failure plot.

Extension: Modify learning rate – how does convergence speed change?

Lesson Summary

Problem Solved: We understand the building block of neural networks and its limitations.

Key Takeaways:

- Perceptron: weighted sum + activation = binary output
- Only learns linearly separable patterns
- XOR problem: needs hidden layers to solve
- Foundation for multi-layer perceptrons (MLPs)

Next Lesson: MLPs and Activations (L34) – adding hidden layers

Memory: Perceptron = single neuron = linear boundary. XOR needs hidden layers.