

Lesson 35: Backpropagation

Data Science with Python – BSc Course

45 Minutes

Learning Objectives

The Problem: Neural networks have millions of weights. How does the network figure out which weights to adjust and by how much?

After this lesson, you will be able to:


- Understand gradient descent optimization
- Interpret loss curves and diagnose training issues
- Configure learning rate and its effects
- Monitor training progress with validation metrics

Finance Application: Training predictive models on financial data

Forward Pass

Computing Predictions


- Input flows through network layer by layer
- Each layer: $z = Wx + b$, then $a = \sigma(z)$



Forward pass: input → hidden(s) → output = prediction

Walking Downhill on the Loss Surface


- Compute gradient: $\frac{\partial L}{\partial w}$ tells direction of steepest increase
- Update: $w_{new} = w_{old} - \eta \cdot \frac{\partial L}{\partial w}$



Gradient descent: repeatedly move opposite to gradient until minimum

Efficiently Computing Gradients


- Chain rule: propagate error backward through layers
- Each layer's gradient depends on layers after it



Backprop = efficient gradient computation via chain rule (not a learning rule)

What Are We Minimizing?


- MSE: regression problems
- Cross-entropy: classification (binary or categorical)



Match loss to problem: MSE for regression, cross-entropy for classification

The Most Important Hyperparameter

- Too high: oscillates, may diverge
- Too low: converges very slowly



Start with 0.001 (Adam default). Reduce if loss is unstable.

Training Curves

Diagnosing Training


- Plot loss vs epochs for train and validation
- Train \downarrow , val \downarrow : good. Train \downarrow , val \uparrow : overfitting



Always monitor validation loss – it shows real generalization

How Much Data Per Update?

- Batch: all data (slow, stable)
- Mini-batch: 32-256 samples (fast, noisy but works)
- Stochastic: 1 sample (very noisy)




Default: mini-batch of 32. Larger batches need larger learning rates.

Volatility Prediction

Finance Application

- Input: lagged returns, volume, past volatility
- Output: next-day volatility (regression)



Monitor training curves carefully – financial data is noisy

Hands-On Exercise (25 min)

Task: Visualize Gradient Descent

- ① Create simple 2D function (e.g., parabola) and plot surface
- ② Implement gradient descent from scratch
- ③ Plot optimization path for different learning rates
- ④ Train Keras model and plot train/val loss curves
- ⑤ Try batch sizes 16, 64, 256 – compare convergence

Deliverable: 3D surface with optimization paths + loss curve comparison.

Extension: Implement momentum and compare to vanilla gradient descent

Lesson Summary

Problem Solved: We understand how neural networks learn through backpropagation and gradient descent.

Key Takeaways:

- Forward pass computes prediction, backward pass computes gradients
- Gradient descent updates weights to minimize loss
- Learning rate controls step size – crucial hyperparameter
- Monitor train/val loss curves to diagnose training

Next Lesson: Overfitting Prevention (L36) – regularization for neural networks

Memory: Backprop = chain rule. Learning rate = step size. Watch val loss.