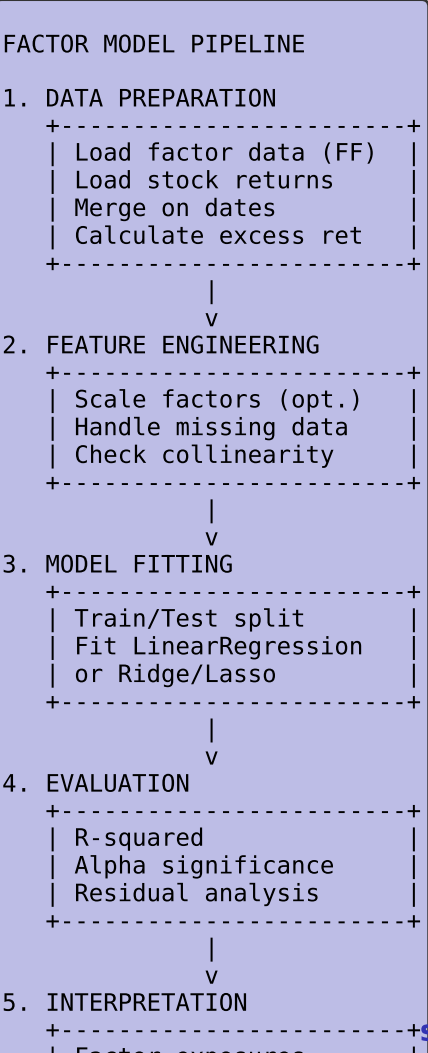


Complete Factor Model Pipeline with sklearn

Pipeline Overview



sklearn Pipeline Code

```
Full sklearn Pipeline Code

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, TimeSeriesSplit
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score, mean_squared_error

# 1. Load and prepare data
stock_ret = pd.read_csv('stock_returns.csv', index_col='date')
ff_factors = pd.read_csv('ff_factors.csv', index_col='date')
data = pd.merge(stock_ret, ff_factors, left_index=True, right_index=True)

# 2. Define features and target
X = data[['Mkt-RF', 'SMB', 'HML', 'MOM']]
y = data['stock_ret'] - data['RF'] # Excess return

# 3. Time series split (NOT random!)
tscv = TimeSeriesSplit(n_splits=5)

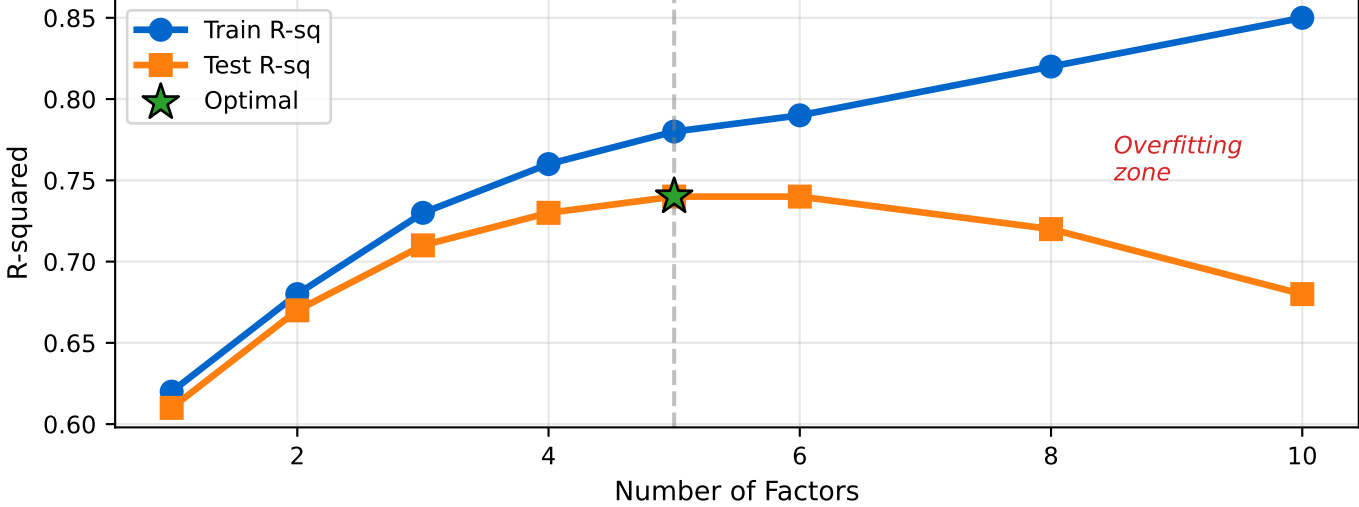
# 4. Create pipeline
pipeline = Pipeline([
    ('scaler', StandardScaler()), # Optional
    ('regressor', Ridge(alpha=0.1))
])

# 5. Cross-validate
scores = []
for train_idx, test_idx in tscv.split(X):
    X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
    y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]

    pipeline.fit(X_train, y_train)
    score = r2_score(y_test, pipeline.predict(X_test))
    scores.append(score)

print(f"CV R-sq: {np.mean(scores):.4f} +/- {np.std(scores):.4f}")
```

Model Complexity: Train vs Test R-squared



Time Series Cross-Validation Results

